

Predict Performance of Processes in Operating Systems Using Machine Learning Techniques

Jiaqian Huang

Department of Computer Science and Engineering
the Ohio State University
Columbus, Ohio, United States
huang.2366@buckeyemail.osu.edu

Abstract—This work presents a comprehensive study on the use of machine learning techniques to predict process performance in operating systems. The proposed approach employs a range of process attributes, such as CPU usage, memory usage, and I/O operations, as input features to train and evaluate machine learning models. Various machine learning algorithms, including Linear Regression and Random Forest Regression, are tested to identify the most effective approach for predicting process performance. The performance of approach is evaluated using data collected from both Linux and macOS operating systems on a local machine. The results demonstrate that some machine learning-based methods can accurately predict process performance. By implementing and evaluating the machine learning models, this study highlights the potential of machine learning in improving the performance and efficiency of operating systems.

Index Terms—Machine learning, Process Performance, Operating System

I. INTRODUCTION

The performance of processes in an operating system is a crucial factor that determines the overall efficiency and responsiveness of the system. Processes are tasks or programs executed by operating system, and their performance can have a significant impact on operating system's overall performance. Efficient allocation of resources, such as CPU time, memory, and I/O operations, is necessary to ensure that each process runs smoothly without delays or crashes.

Predicting process performance provides numerous benefits, including early detection of performance issues, improved resource utilization, better capacity planning, reduced downtime, enhanced user experience, and cost savings.

Users on any operating system can check process information in the terminal using built-in command lines. This information typically includes unique process ID for each task, username of the task owner, task priority, total virtual memory used, actual physical memory consumption, shared memory size in kilobytes, task state, CPU and memory usage, CPU time, and command being run. By leveraging large amounts of data with multiple attributes, machine learning algorithms can identify patterns and make accurate predictions about process performance, which motivates researchers to design and develop possible predictive models.

In this work, the application of predictive machine learning models to estimate the performance of processes is investigated. This work demonstrates that the data collected on a

local machine can be generalized to represent the performance of processes on machines with different hardware specifications. Through the experiments, Random Forest Regression is efficient in estimating the performance by using the process information on both Linux and macOS operating systems. But Linear Regression is not as efficient as Random Forest Regression. Furthermore, this study concludes that the performance of processes can be accurately predicted using data collected by any user on both Linux and macOS operating systems.

This paper is organized as follows. Section II provides an overview of related work in predicting performance. Section III describes the study's design, including methods used to collect data on both Linux and macOS operating systems and predictive machine learning models employed. In Section IV, a detailed evaluation of experiments is presented, including setup and results. Section V discusses the challenges encountered during the experiments. Finally, in Section VI, the findings are summarized, and a conclusion is provided for this work.

II. RELATED WORK

Machine learning algorithms have gained popularity in accurately predicting performance and behavior in computer science. Previous works, such as Fu, Gupta, Mittal and Ratnasam [8], applied machine learning to predict application and use-case performance, but the prediction error remained high for realistic scenarios. Other studies have applied machine learning algorithms to predict computer operating system security, including whether a newly discovered vulnerability would enable attackers to cause denial of service [3] and developing a light-weight machine learning engine in the kernel space component [6]. Aslam, Sarwar and Batool [7] focused on improving CPU scheduling by applying Bayesian Decision Theory. However, there is a need for more general models that can accurately predict process performance across different operating systems and user scenarios. This study aims to address this gap by developing predictive machine learning models for both Linux and macOS operating systems using data generated by local machine from diverse users.

III. DESIGN

Here's an overview of each step in the design:

- **Data collection:** This step involves gathering process-related data from both Linux and macOS operating systems. The data are used as input features for training and evaluating machine learning models.
- **Data preprocessing:** After collecting the data, it is necessary to clean and preprocess them to ensure that the dataset is consistent and ready for use in machine learning models. This involves removing any duplicates or incomplete data, normalizing the data to a common scale, and encoding categorical variables.
- **Machine learning model building:** In this step, several machine learning models are developed and trained by using the preprocessed data. The goal is to identify the most suitable approach for predicting process performance. Various machine learning algorithms are experimented, including Linear Regression and Random Forest Regression.
- **Evaluation:** To evaluate the accuracy of machine learning models, couple evaluation metrics are applied on models. This will help verify the performance of models and identify any potential issues.

A. Data Collecting

Data collecting is an essential part in this project. Process-related data are required to be collected from various operating systems.

To collect the necessary data, a combination of built-in system utilities and custom scripts is applied in this work, which allows users to collect process attributes such as CPU usage, memory usage, I/O operations, and response time.

To ensure the accuracy and comprehension of data, collecting data from multiple operating systems at various intervals takes the next step. It is also recommended to cover a variety of process-related data to represent different behaviors and collect them at different hours of a day such as peak hours and off-peak hours.

Specifically, Linux operating system and macOS operating system are chosen and focused on in this work, as they are widely used in real-world settings and will provide representative data.

1) Linux Operating System

A kernel module, also referred to as a driver, is a software component that can be dynamically loaded or unloaded into the kernel of an operating system. It has the capability to interact with hardware components and offer services to user-space applications. The kernel module is coded in the C language, a low-level programming language that conforms to the kernel's programming interface (API) and driver model.

In this project, a kernel module is implemented to gather data generated by operating system and store it in a virtual file on Linux operating system. Additionally, a C file is executed in user-level to export contents of the virtual file into an actual file.

The motivation behind creating a kernel module to collect process-related data on Linux operating system is that it can be dynamically loaded and unloaded without requiring a

reboot of operating system. This flexibility in managing system resources does not impose a heavy burden on the kernel of Linux operating system to handle.

2) macOS Operating System

While it is technically possible to create a kernel module on macOS, it is generally not recommended due to the tightly controlled kernel architecture, the proprietary nature of the kernel code, and the availability of high-level programming tools and libraries for creating user-space applications. As a result, the kernel module is not loaded on macOS. Instead, one Python script is developed to gather data on macOS operating system.

One advantage of using Python in a macOS environment is that it is relatively simple because macOS operating system includes pre-installed development tools like a C compiler and the XCode development environment. These tools are critical for constructing Python from source code, and their availability on macOS makes it effortless to set up and use Python on a MacBook.

B. Data Preprocessing

After collecting process-related data from both Linux and macOS operating systems, it is essential to conduct a comprehensive cleaning and preprocessing of the data to ensure consistency. This involves removing any irrelevant data that is not related to processes. Handling missing data is another important step, which can be achieved by imputing values based on the mean or median of the respective attribute. Finally, techniques like one-hot encoding or label encoding can be used to encode categorical variables such as process state in dataset into numerical values, allowing machine learning models to process the data more effectively.

By following these data cleaning and preprocessing steps, machine learning models can be trained on high-quality and consistent data. These processes improve models' accuracy and performance in predicting process performance on operating systems.

C. Predictive Models

In this project, data collected from both Linux and macOS operating systems are used to build predictive models to estimate the performance of processes in operating systems respectively. 2 types of models are developed: Linear Regression and Random Forest Regression.

1) Linear Regression

Linear Regression is one of the simplest and most commonly used linear models. The primary goal of Linear Regression is to identify a linear equation that closely fits the data and best describes the relationship between predictor variable and response variable [4]. The predictor variable is also named as independent variable [4]. The response variable is also named as dependent variable and researchers apply different ways to predict this variable [4]. Linear Regression allows for identification of the most important predictor variables, estimation of magnitude and direction of their effects, and

avoidance of overfitting by regularizing the estimates [4]. The multiple Linear Regression equation is as follows [4]:

$$y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n + e \quad (1)$$

where y is dependent variable or predicted outcome variable, and x_1, x_2, \dots, x_n are independent variables or predictors, and b_0 is intercept or expected value of y when all independent variables are equal to 0, and b_1, b_2, \dots, b_n are regression coefficients which contributes the effect of each independent variable on y , and e is error term which represents the part of y that is not explained by independent variables [4].

In this study, the scikit-learn (sklearn) Python library is utilized to apply Linear Regression. The scikit-learn (sklearn) Python library provides a comprehensive set of tools for data preprocessing, model selection, and evaluation. Those tools are utilized to build models that predict the target variable based on a set of input features.

2) Random Forest Regression

Random Forest Regression is an extension of the decision tree algorithm that builds multiple decision trees and merges them together to improve accuracy and reduce overfitting [1]. The Random Forest Regression works by creating a random sample of training data and then building decision trees on each sample [1]. The algorithm then combines the results of all decision trees to make predictions, reducing variance and improving prediction accuracy [1].

One of the key advantages of Random Forest Regression is its ability to handle high-dimensional data and capture complex nonlinear relationships between input and target variables [1]. It is also robust to outliers and missing data, making it a versatile choice for a wide range of regression tasks [2]. Besides, variable importance can tell developers how much the Random Forest Regression model uses each variable to make accurate predictions [1].

In this study, Random Forest Regression was applied using the scikit-learn (sklearn) Python library. Overall, the algorithm is proved to be a powerful tool in predicting process performance, particularly when dealing with complex datasets.

3) Evaluation and Validation

Some machine learning models do not provide an estimated predictive error for themselves. Therefore, it is necessary to introduce an evaluation metric to ensure that the performance of a predictive model provides accurate results. Typically, an evaluation metric requires training the model on a training dataset, using the model to make predictions on a holdout dataset which is also referred as testing dataset, and then comparing the predicted values with the expected values.

In this work, R-squared [4] is applied as one of evaluation metrics. R-squared, also known as the coefficient of determination, measures whether a prediction for a regression model is efficient or not and yields a score between 0 and 1 [1]. A value of 1 corresponds to a perfect prediction, and a value of 0 corresponds to a constant model that just predicts the mean of the training set responses [1]. Thus, a model with relatively good model fit may have a high R-squared score value. In

contrast, if the model does not fit the data, it may get a low R-squared score value.

The form of it is Equation [4]:

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} \quad (2)$$

where SS_{res} is the sum of squares of the residuals, and SS_{tot} is the total sum of squares [4].

However, it is important to note that R-squared should not be used as the sole criterion for evaluating a regression model. Considering other metrics, such as Mean Squared Error (MSE) or cross validation, can ensure the model is accurately predicting the outcome variable and not simply overfitting the data.

Thus, Mean Squared Error (MSE) [4] is applied as well. Mean Squared Error (MSE) is the sum of the squared differences between the predictions and the true values [1]. The following equation shows the formula for calculating Mean Squared Error (MSE) [4]:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (3)$$

where n is the number of observations in the dataset, and y_i is the actual value of the dependent variable for observation i , and \hat{y}_i is the predicted value of the dependent variable for observation i [4].

While Mean Squared Error (MSE) is applied, it should be taken into account that Mean Squared Error (MSE) can be influenced by outliers, known as extreme data, in the dataset. Therefore, the combination of Mean Squared Error (MSE), R-squared and cross validation is persuasive to obtain a comprehensive evaluation of a model performance.

Cross validation is utilized in this work too, which is a commonly used statistical method for evaluating the generalization performance of machine learning models [1]. It is considered more stable and thorough than simply splitting the data into training and test sets [1]. In the scikit-learn (sklearn) Python library, the `cross_val_score` function can be used to perform cross validation [1]. The most commonly used version of cross validation is k-fold cross-validation, where the data is divided into k parts or folds of approximately equal size [1]. For each of the k folds, a sequence of models is trained and evaluated [1]. The first model is trained using the first fold as the test set and the remaining folds as the training set [1]. The accuracy is then evaluated on the first fold [1]. The process is repeated using a different fold as the test set each time until all the k folds have been used as test sets [1]. The accuracy scores obtained from all the k folds are averaged to give an estimate of the model's generalization performance [1]. The scikit-learn (sklearn) Python library supports several types of cross validation strategies. By employing cross validation, machine learning models can be evaluated on a limited sample of data in a more robust and reliable manner.

In addition, model validation for Linear Regression model can be considered. Linear Regression needs to match couple

assumptions. For instance, Linear Regression model is based on the assumption that there exists a linear relationship between predictors and response variable [5]. If this assumption is not met, the accuracy of predictions may be significantly compromised, and the validity of conclusions drawn from the model may be questionable [5].

An additional key assumption in linear regression is that the error terms have a constant variance, denoted as $\text{Var}(\epsilon_i) = \sigma^2$ [5]. This assumption is essential for valid inferences and prediction intervals based on the model [5]. Unfortunately, in some cases, the variances of error terms may be non-constant, resulting in heteroscedasticity [5]. For example, the variances of error terms may increase or decrease with the value of response variable [5]. One way to detect heteroscedasticity is by examining residual plots, which may exhibit a funnel shape or some other systematic pattern [5]. If heteroscedasticity is detected, one may need to consider alternative regression methods, such as weighted least squares regression, to account for non-constant variance of errors [5].

Besides, multicollinearity is a phenomenon where three or more variables are linearly related, even if no pair of variables has a particularly high correlation [5]. Variance Inflation Factor (VIF) is a measure that can help assess the extent of multicollinearity [5]. A Variance Inflation Factor (VIF) value of 1 indicates the complete absence of collinearity, while a value greater than 1 suggests the presence of some degree of collinearity among predictors [5]. In practice, a Variance Inflation Factor (VIF) value below 5 is considered acceptable and does not require further investigation to correct multicollinearity [5]. However, a Variance Inflation Factor (VIF) value exceeding 5 or 10 indicates a problematic amount of collinearity that needs to be addressed [5]. Variance Inflation Factor (VIF) equation is given by [4]:

$$VIF_i = \frac{1}{1 - R_i^2} \quad (4)$$

where R_i^2 is the coefficient of determination mentioned in (2) that regresses the i th predictor variable on all other predictor variables [4].

IV. EVALUATION

A. Experiment Setup

The experimental setup involves collecting the data, preprocessing the dataset, implementing predictive models, evaluating and validating their performance using multiple metrics, and measuring and comparing their computational efficiency.

1) Data Collecting and Input Values on Both Operating Systems

Before installing kernel module, `proclog.c`, a file built for the kernel module, was compiled. It contained several kernel functions to facilitate the logging of information related to processes into a virtual file. During the compilation process, several files with the same name but different file types were generated.

To start logging data, the file with the `.ko` extension was loaded on Linux operating system. The load process was

simple by executing the make file command containing custom actions. Then this custom kernel module was added to the list of running kernel modules, and no additional steps were required to collect process-related information. When the `.ko` file was loaded successfully, the user was reminded via a message. If some issues occurred after installation, the user was able to execute the make file command to unload the kernel module.

The kernel module created a virtual file stored in the `"/proc/"` folder. The virtual file can be written to by kernel-level functions but not by user-level applications. As Linux environment was active after installation, the virtual file was consistently updated by the kernel module to log process-related data. The file grew in size dramatically in a very short time. However, due to its attribute as a virtual file, it did not take up significantly large space on hard disk.

Since kernel modules are usually not designed to write actual files, as they operate at a low level in the operating system and are mainly used to interact with hardware components, system resources, and other kernel components, a user-level C file was designed to write the contents of the virtual file into an actual file. Furthermore, because it is not possible to export comprehensive data from virtual data due to its continuous update by the kernel module for data logging, an input step for the user to manually input a positive integer was created to set up the number of lines to export. The kernel module operated while the system was on. When the file was open for exporting, the read never reached the end, and the program for data export never terminated. In this experiment, the input number was 3000 which represented the user requested 3000 records of process-related information from virtual file.

On macOS Operating System, before the execution of Python program built for data collecting, it was required to install Python-related packages for program to import and use. In the program, a library called `psutil` was applied and it retrieved process-related data with a few lines of code in Python program. The program set up an input step for user to choose an interval and end time for data collection. Python program retrieved process data every interval that the user input and stored the data in a file in csv format. The entire process for data collection finished within the user-specified time. In this experiment, the input interval was set as 5 seconds and the time for data collection finished in 30 seconds.

Any user is capable of generating own dataset containing process-related information on either Linux operating system or macOS operating system, provided that the user has access to the relevant operating system. Since the performance of processes may vary between different computers, the generated data may also differ. The dataset used in the experiment was collected on a MacBook and Linux operating system set up in a virtual machine.

2) Data Preprocessing and Input Values

In this experiment, the dataset was subjected to several preprocessing steps to make it suitable for machine learning models.

For Linux operating system, preprocessing was performed

during the export process to ensure compatibility with common data analysis tools. Each row was thoroughly checked, and all white spaces were replaced with commas to match the CSV file format. This eliminated the need for an additional TXT to CSV conversion process. Conversely, data collection on macOS operating systems was streamlined using the csv Python library. This library enabled writing the header of a CSV file and directly appending records into rows, requiring no additional modifications to prepare the dataset in CSV format.

To prepare the dataset for machine learning models, several preprocessing steps were taken. Firstly, the dataset was checked for missing data, and unnecessary data were dropped, such as columns containing the same data for all rows. Next, label encoding was used to convert categorical data to numerical data, since one-hot encoding led to high memory consumption due to the large number of categories. Besides, all rows that had outliers in at least one column were removed.

To enhance the user experience, a feature was implemented where the program prompted the user to input the data to process and the type of model to build upon execution. This approach reduced the wait time for the user, as the model building process took several minutes or longer. By allowing the user to specify desired model before the processing began, the program was able to let user wait for one model to be generated.

In the process to build the machine learning model, the dataset was split into training and testing sets using a 3:1 ratio and a random seed of 42 was set, which is a commonly used integer for random seeds. Setting the random seed ensures that other developers who use the same dataset and source code can reproduce the exact output. The training set was used to train the predictive models, while the testing set was used to evaluate their performance.

Overall, these preprocessing steps ensured that the dataset was of high quality and suitable for use in machine learning models.

3) Predictive Models Implementing

Since the accesses to process-related data on different operating systems were different, the predicted columns in data frame were different based on data for different platforms.

If the user chose Linear Regression for process prediction, in preparation for building Linear Regression model, several specific steps were taken to ensure that the data were appropriate for this type of analysis. To scale and center the features of a dataset, standard scaler was applied and made data prepared in a standardized distribution of features, ensuring that each feature contributed equally to the learning process and the impact of outliers was reduced. Next, the linear relationship was checked by scatter plotting “Actual value Vs Predicted value” and multivariate normality was checked with quantile-quantile (q-q) plot. “Residual value Vs Fitted value” scatter plot was checked too. Another step was to check for multicollinearity. To address this issue, highly correlated independent variables were identified and removed using the Variance Inflation Factors (VIF) technique. A check




Fig. 1. Machine Learning Algorithms applying Linux data.




Fig. 2. Machine Learning Algorithms applying MacOS data.

function was used to calculate the Variance Inflation Factors (VIF), ensuring that the data had low to moderate levels of multicollinearity. Linear Regression assumes that there is minimal or no multicollinearity in the dataset, so these steps were crucial to ensure the validity of the models.

If choosing Random Forest Regression, the user needed to expect a longer processing time because there were several parameters to set up, and GridSearchCV function from the scikit-learn (sklearn) Python library required time to calculate and select the best parameters from various options. The GridSearchCV function was utilized to select the optimal hyperparameters for Random Forest Regression model on Linux and MacOS data. By utilizing these parameters, the model was able to achieve optimal performance for the given datasets. And variable importance was calculated.

B. Experiment Result

For Random Forest Regression, the selected parameters for model based on Linux data were that *maxDepth* was 7, and *maxFeatures* was *sqrt*, and *minSamplesLeaf* was 1, and *nEstimators* was 500, and *randomstate* was 42. Similarly, based on MacOS data, GridSearchCV function selected parameters that *maxDepth* was 7, and *maxFeatures* was *sqrt*, and *minSamplesLeaf* was 1, and *nEstimators* was 300, and *randomstate* was 2017.

After analyzing Fig. 1 and Fig. 2, it is evident that Linear Regression was not an effective machine learning algorithm for neither Linux dataset nor macOS dataset. The data exhibited insufficient linearity, which resulted in inaccurate predictions.

TABLE I
EVALUATION METRICS CALCULATED ON LINUX DATA

Methods	Linear Regression	Random Forest Regression
Test MSE	8.805821e+15	3.226405e+14
Test RMSE	9.383933e+07	1.796220e+07
Test R2	0.173071	0.99298
MSE Mean	7.684483e+15	3.127584e+14
RMSE Mean	8.733552e+07	1.760835e+07

TABLE II
EVALUATION METRICS CALCULATED ON MACOS DATA

Methods	Linear Regression	Random Forest Regression
Test MSE	3.895207e+11	6.636878e+10
Test RMSE	624115.918959	257621.382557
Test R2	0.514342	0.992468
MSE Mean	3.346303e+11	9.660352e+10
RMSE Mean	576119.614839	302881.154619

Conversely, Random Forest Regression algorithm proved to be successful in fitting the data and produced better results.

Based on Table I and Table II, Random Forest Regression outperformed the other models in terms of both training performance and testing performance while data collected on MacBook and Linux operating system were processed respectively. It achieved the lower MSE, lower RMSE and higher R-squared scores on testing data, indicating a higher level of accuracy in its predictions and a better fit to the data. In contrast, Linear Regression exhibited a lower level of accuracy in its predictions, with higher MSE, higher RMSE and lower R-squared scores on testing data.

Overall, Random Forest Regression may be a better choice for predicting the outcome variable based on the given input variables. However, further investigation and evaluation may be necessary to determine the optimal model for a particular application.

V. CHALLENGES AND FUTURE WORK

The development of kernel modules poses various challenges that require careful consideration. Kernel module development presents a significant debugging challenge. Unlike regular user space applications, kernel modules lack access to detailed debugging information provided by operating system, making it arduous for developers to identify and resolve issues that may arise during the module's development process. Developers need to rely on various strategies to identify errors in their code, such as kernel logs, print statements, and kernel debuggers. However, these methods can be time-consuming and inefficient, especially when dealing with complex kernel modules. Despite this challenge, the use of virtual machines, which is the only method applied in this work and make the work relatively easy, can help isolate and reproduce kernel module-related problems in a controlled environment, facilitating the debugging process. However, it is not efficient enough for developers to deal with the tedious process of debugging in this work and the potential issues may occur without being caught by operating system or developers.

Furthermore, not all process-related information is accessible on both operating systems. Some inherent limitations exist within kernel modules to access information, resulting in incomplete data collection by the kernel module. In addition, built-in Python functions are not capable to collect data related to io counter and net counter for each process, which are critical factors for process performance.

In future work for topic of process performance, other models with potentially superior performance can also be explored. Additionally, greater attention can be paid to model validation. To validate a model, checking certain assumptions during data preparation is necessary, as is the case with Random Forest Regression, and it is recommended to perform model validation for this type of model as well. Furthermore, more investigation of variable importance can be implemented to select the useful variable to interpret the most fitted forest [2]. Besides, the time required to load the machine learning Python script is considerable and takes several minutes. Therefore, it is important to explore ways to shorten the execution time and improve the efficiency of process to build predictive models.

VI. CONCLUSION

In conclusion, this work explored the use of Linear Regression and Random Forest Regression models to estimate process performance and behavior in Linux and macOS operating systems. The study showed that only Random Forest Regression can accurately predict process performance based on collected data from both operating systems and it outperformed the Linear Regression in terms of training and testing accuracy. These results demonstrate that not all machine learning algorithms are suitable for prediction and there is a potential of machine learning techniques in analyzing and understanding process performance in different operating systems. These predictions made by Random Forest Regression can provide valuable insights into how processes are performing, which predicting process performance can help achieve a more efficient system performance and enable programs to run efficiently, which is essential for optimal system functionality. The developed models provide a foundation for future research in this field and may lead to improved system optimization and performance.

REFERENCES

- [1] A. C. Müller and S. Guido, "Introduction to Machine Learning with Python," O'Reilly Media, Inc., Sept 2016
- [2] A. Cutler, D. R. Cutler and J. R. Stevens, "Random Forests," Ensemble Machine Learning: Methods and Applications, pp. 157-176, January 2011
- [3] F. Alenezi, and C. P. Tsokos, "Machine Learning Approach to Predict Computer Operating Systems Vulnerabilities," 2020 3rd International Conference on Computer Applications and Information Security (ICCAIS), pp. 1 – 6, March 2020
- [4] G. A. F. Seber and Alan J. Lee, "Linear Regression Analysis", Wiley, Febr 2003
- [5] G. James, D. Witten, T. Hastie and R. Tibshirani, "An Introduction to Statistical Learning with Applications in R," Springer, Jun 2013
- [6] I. U. Akgun, A. S. Aydin and E. Zadok, "KMLib: Towards Machine Learning For Operating Systems," On-device Intelligence Workshop MLSys 2020, Febr 2020

- [7] N. Aslam, N. Sarwar and A. Batool, "Designing a Model for improving CPU Scheduling by using Machine Learning," (IJCSIS) International Journal of Computer Science and Information Security, Vol. 14, No. 10, Oct 2016
- [8] S. Fu, S. Gupta, R. Mittal and S. Ratnasamy, "On the Use of ML for Blackbox System Performance Prediction," Symposium on Networked Systems Design and Implementation, Apr 2021